

A2-02-Image-Segmentation

June 10, 2026

1 A2-02 · Image Segmentation — do skip connections matter?

This notebook is fully self-contained. The U-Net/ResNet-18 model, the Oxford-Pet dataset and the training/evaluation loop are defined in the library cells below (inlined verbatim from the `segmentation/` package), and the experiment cells **train both variants and evaluate them** here — the same experiment as `run.py --model unet_resnet18[...] --train`.

Two **U-Net** decoders share the *same* ImageNet-pretrained **ResNet-18** encoder; the only difference is whether each decoder stage concatenates the matching encoder feature map (**skip connections**). Training both on Oxford-IIIT Pet for 20 epochs isolates exactly what skip connections contribute.

Re-runs are idempotent — cached checkpoints under `checkpoints/` are reused unless `RETRAIN=True`. The Oxford-Pet dataset downloads on first use.

```
[1]: import os, sys, json, time, math, glob, random
# Run from the repository root so cfg/ weights/ data/ asset paths resolve.
if os.path.basename(os.getcwd()) == 'notebooks':
    os.chdir('..')
import numpy as np
import pandas as pd
import torch
import matplotlib.pyplot as plt
%matplotlib inline
from IPython.display import Image as IPyImage, display

def set_seed(seed=42):
    """Identical to run.py's set_seed - reseed right before each train so a
    fresh notebook run and run.py consume the RNG in the same order and produce
    the same numbers (exact for the num_workers=0 path)."""
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

def record_metric(key, value, path='results/metrics.json'):
    """Persist a scalar to results/metrics.json - identical helper and keys to
    run.py, so notebook and CLI runs populate the same metrics file."""
```

```

os.makedirs(os.path.dirname(path), exist_ok=True)
data = {}
if os.path.exists(path):
    try:
        data = json.load(open(path))
    except json.JSONDecodeError:
        data = {}
data[key] = value
json.dump(data, open(path, 'w'), indent=2)

def show_if(path, width=None):
    """Display a static diagram if present (optional documentation asset)."""
    if os.path.exists(path):
        display(IPyImage(path, width=width))
    else:
        print(f'[optional diagram not found: {path}] - see lab_note/img')

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print('device:', device, '|',
      torch.cuda.get_device_name(0) if torch.cuda.is_available() else 'CPU')

```

device: cuda | NVIDIA GeForce RTX 5060 Ti

1.1 0. Experiment configuration

Defaults match the assignment (20 epochs per variant). Lower SEG_EPOCHS for a quick smoke test, or set RETRAIN=True to ignore cached checkpoints.

```

[2]: SEG_EPOCHS = 20      # assignment spec: 20 epochs per variant
BATCH_SIZE = 16
IMG_SIZE = 128
LR = 1e-3
NUM_WORKERS = 0      # keep 0 for exact determinism (matches run.py defaults)
SEED = 42            # same seed as run.py --seed; reseeded before each train
RETRAIN = True       # True -> ignore cached checkpoints/ and train from scratch

```

1.2 1. Library code (inlined from the segmentation/ package)

The model, dataset and training/eval loop. Run these once, top to bottom.

1a. Model — UNetResNet18 with a use_skip switch (build_segmentation_model picks the variant).

```

[3]: # =====
# Inlined verbatim from segmentation/models.py (self-contained notebook)
# =====
"""U-Net with a pretrained ResNet-18 encoder - skip vs no-skip ablation.

Both variants share the identical ImageNet ResNet-18 encoder and the same

```

decoder depth/width; the only difference is whether each decoder stage concatenates the matching encoder feature map (`use_skip`). That isolates the contribution of skip connections (Exercise 1).

"""

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision.models as models
```

```
class DoubleConv(nn.Module):
```

"""Conv → BN → ReLU, twice - the core U-Net block."""

```
    def __init__(self, in_ch, out_ch):
        super().__init__()
        self.block = nn.Sequential(
            nn.Conv2d(in_ch, out_ch, 3, padding=1, bias=False),
            nn.BatchNorm2d(out_ch), nn.ReLU(inplace=True),
            nn.Conv2d(out_ch, out_ch, 3, padding=1, bias=False),
            nn.BatchNorm2d(out_ch), nn.ReLU(inplace=True),
        )
```

```
    def forward(self, x):
        return self.block(x)
```

```
class UNetResNet18(nn.Module):
```

"""U-Net / ResNet-18. `use_skip=False` removes every encoder→decoder_↪skip.

Encoder feature maps on a 128×128 input:

stem_conv → H/2, 64ch (s0)

stem_pool → H/4

layer1 → H/4, 64ch (s1)

layer2 → H/8, 128ch (s2)

layer3 → H/16, 256ch (s3)

layer4 → H/32, 512ch (s4) → bottleneck

Decoder upsamples H/32→H over 5 stages.

"""

```
    def __init__(self, n_classes=3, pretrained=True, use_skip=True):
        super().__init__()
        self.use_skip = use_skip
```

```
        weights = "IMAGENET1K_V1" if pretrained else None
```

```
        resnet = models.resnet18(weights=weights)
```

```

        self.stem_conv = nn.Sequential(resnet.conv1, resnet.bn1, resnet.relu)
↪ # H/2, 64
        self.stem_pool = resnet.maxpool
↪ # H/4
        self.enc1 = resnet.layer1 # H/4, 64
        self.enc2 = resnet.layer2 # H/8, 128
        self.enc3 = resnet.layer3 # H/16, 256
        self.enc4 = resnet.layer4 # H/32, 512

        self.bottleneck = DoubleConv(512, 1024)

        # skip-channel count added to each decoder stage (0 when skips disabled)
        s4, s3, s2, s1, s0 = (512, 256, 128, 64, 64) if use_skip else (0, 0, 0,
↪ 0, 0)

        self.up4 = nn.ConvTranspose2d(1024, 512, 2, stride=2)
        self.dec4 = DoubleConv(512 + s4, 512)
        self.up3 = nn.ConvTranspose2d(512, 256, 2, stride=2)
        self.dec3 = DoubleConv(256 + s3, 256)
        self.up2 = nn.ConvTranspose2d(256, 128, 2, stride=2)
        self.dec2 = DoubleConv(128 + s2, 128)
        self.up1 = nn.ConvTranspose2d(128, 64, 2, stride=2)
        self.dec1 = DoubleConv(64 + s1, 64)
        self.up0 = nn.ConvTranspose2d(64, 32, 2, stride=2)
        self.dec0 = DoubleConv(32 + s0, 32)

        self.output = nn.Conv2d(32, n_classes, kernel_size=1)

    def _join(self, x, skip):
        """Concatenate skip (if enabled), upsampling skip to match x if needed.
↪ """
        if not self.use_skip:
            return x
        if x.shape[2:] != skip.shape[2:]:
            skip = F.interpolate(skip, size=x.shape[2:])
        return torch.cat([skip, x], dim=1)

    def forward(self, x):
        s0 = self.stem_conv(x)
        sp = self.stem_pool(s0)
        s1 = self.enc1(sp)
        s2 = self.enc2(s1)
        s3 = self.enc3(s2)
        s4 = self.enc4(s3)

        x = self.bottleneck(s4)

```

```

        x = self.dec4(self._join(self.up4(x), s4))
        x = self.dec3(self._join(self.up3(x), s3))
        x = self.dec2(self._join(self.up2(x), s2))
        x = self.dec1(self._join(self.up1(x), s1))
        x = self.dec0(self._join(self.up0(x), s0))
        return self.output(x)

def build_segmentation_model(name, n_classes=3, pretrained=True):
    """``unet_resnet18`` → skips on; ``unet_resnet18_no_skip`` → skips off."""
    if name == "unet_resnet18":
        return UNetResNet18(n_classes, pretrained, use_skip=True)
    if name == "unet_resnet18_no_skip":
        return UNetResNet18(n_classes, pretrained, use_skip=False)
    raise ValueError(f"Unknown segmentation model: {name!r}")

```

1b. Dataset — Oxford-IIIT Pet (3 classes: pet / background / border), with denormalize for display.

```

[4]: # =====
# Inlined verbatim from segmentation/dataset.py (self-contained notebook)
# =====
"""Oxford-IIIT Pet segmentation dataset (3 classes: pet / background / border).
↪ """

import numpy as np
import torch
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms
from torchvision.datasets import OxfordIIITPet

CLASS_NAMES = ["Pet", "Background", "Border"]
CLASS_COLORS = np.array([[255, 100, 100], [100, 100, 255], [255, 255, 100]],
↪ dtype=np.uint8)
IMAGENET_MEAN = torch.tensor([0.485, 0.456, 0.406]).view(3, 1, 1)
IMAGENET_STD = torch.tensor([0.229, 0.224, 0.225]).view(3, 1, 1)

class PetSegDataset(Dataset):
    def __init__(self, base, size=128):
        self.ds = base
        self.img_tf = transforms.Compose([
            transforms.Resize((size, size)),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]),
        ])
        self.mask_tf = transforms.Compose([

```

```

        transforms.Resize((size, size), interpolation=transforms.
↳InterpolationMode.NEAREST),
        transforms.PILToTensor(),
    ])

    def __len__(self):
        return len(self.ds)

    def __getitem__(self, idx):
        img, mask = self.ds[idx]
        img = self.img_tf(img)
        mask = (self.mask_tf(mask).squeeze(0).long() - 1).clamp(0, 2) #↳
↳{1,2,3} → {0,1,2}
        return img, mask

def get_pet_loaders(data_dir="data", size=128, batch_size=16, num_workers=0,↳
↳download=True):
    """Return ``(train_loader, test_loader, train_data, test_data)``."""
    train_raw = OxfordIIITPet(data_dir, split="trainval",↳
↳target_types="segmentation", download=download)
    test_raw = OxfordIIITPet(data_dir, split="test",↳
↳target_types="segmentation", download=download)
    train_data = PetSegDataset(train_raw, size)
    test_data = PetSegDataset(test_raw, size)
    train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True,↳
↳num_workers=num_workers)
    test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False,↳
↳num_workers=num_workers)
    return train_loader, test_loader, train_data, test_data

def denormalize(img):
    """Undo ImageNet normalisation → HWC float image in [0,1] for display."""
    return torch.clamp(img.cpu() * IMAGENET_STD + IMAGENET_MEAN, 0, 1).
↳permute(1, 2, 0).numpy()

```

1c. Training + evaluation — per_class_iou, evaluate_segmentation (mIoU), and train_segmentation.

```

[5]: # =====
# Inlined verbatim from segmentation/train.py (self-contained notebook)
# =====
"""U-Net training + mIoU evaluation for the Oxford-Pet segmentation task."""

import json
import os

```

```

import time

import numpy as np
import torch
import torch.nn as nn
from tqdm.auto import tqdm

def per_class_iou(pred, target, n_classes=3):
    """Per-class IoU for a batch of logits; returns ``list`` (``nan`` if
    ↪absent)."""
    pred = pred.argmax(dim=1)
    ious = []
    for c in range(n_classes):
        inter = ((pred == c) & (target == c)).sum().float()
        union = ((pred == c) | (target == c)).sum().float()
        ious.append((inter / union).item() if union > 0 else float("nan"))
    return ious

def compute_miou(pred, target, n_classes=3):
    ious = [v for v in per_class_iou(pred, target, n_classes) if not np.
    ↪isnan(v)]
    return float(np.mean(ious)) if ious else 0.0

@torch.no_grad()
def evaluate_segmentation(model, loader, device, n_classes=3):
    """Return ``(mIoU, per_class_IoU_list)`` over a loader."""
    model.eval()
    mious, per_cls = [], [[] for _ in range(n_classes)]
    for imgs, masks in loader:
        out = model(imgs.to(device))
        masks = masks.to(device)
        mious.append(compute_miou(out, masks, n_classes))
        for c, v in enumerate(per_class_iou(out, masks, n_classes)):
            if not np.isnan(v):
                per_cls[c].append(v)
    miou = float(np.mean(mious)) if mious else 0.0
    pc = [float(np.mean(v)) if v else 0.0 for v in per_cls]
    return miou, pc

def train_segmentation(model, train_loader, test_loader, device, epochs=20,
    ↪lr=1e-3,
                        ckpt_path="UNET_resnet18_pet.pt", tag="UNET_resnet18",
                        ckpt_dir="checkpoints"):

```

```

    """Train a U-Net variant; save weights + history; return the history dict.
    ↪ """
    os.makedirs(ckpt_dir, exist_ok=True)
    model = model.to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=lr)
    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=5, gamma=0.
    ↪5)

    history = {"loss": [], "miou": [], "time": [], "per_class": [], "tag": tag}
    for epoch in range(epochs):
        model.train()
        t0 = time.time()
        ep_loss = []
        for imgs, masks in tqdm(train_loader, desc=f"[{tag}] epoch {epoch+1}/
    ↪{epochs}"):
            imgs, masks = imgs.to(device), masks.to(device)
            loss = criterion(model(imgs), masks)
            optimizer.zero_grad(); loss.backward(); optimizer.step()
            ep_loss.append(loss.item())
        scheduler.step()
        elapsed = time.time() - t0
        miou, pc = evaluate_segmentation(model, test_loader, device)
        history["loss"].append(float(np.mean(ep_loss)))
        history["miou"].append(miou)
        history["time"].append(elapsed)
        history["per_class"].append(pc)
        print(f"Epoch {epoch+1:02d}/{epochs} | Loss {history['loss'][-1]:.4f} |
    ↪
            f"mIoU {miou:.4f} | {elapsed:.1f}s")

    torch.save(model.state_dict(), ckpt_path)
    torch.save(model.state_dict(), os.path.join(ckpt_dir, f"{tag}.pt"))
    with open(os.path.join(ckpt_dir, f"history_{tag}.json"), "w") as fp:
        json.dump(history, fp, indent=2)
    print(f"Saved → {ckpt_path} and {ckpt_dir}/history_{tag}.json")
    return history


```

1.3 2. Semantic vs instance segmentation

	Semantic	Instance
Question	class of each pixel	which <i>object</i> each pixel belongs to
Same-class objects	merged	separated
Typical model	U-Net / FCN	Mask R-CNN

Mask R-CNN = Faster R-CNN + a per-instance mask head (RoI-Align, 28×28 mask). The course U-Net diagram — grey arrows are the skip connections:

```
[6]: show_if('lab_note/img/unet_arch.png', width=720)
```



A live **Mask R-CNN** (pretrained COCO, torchvision) on the test image — boxes from the Faster-R-CNN head, per-instance masks from the mask head. This is the *instance* counterpart to the *semantic* U-Net we train below:


```
[7]: import cv2
import torchvision.transforms.functional as TF
from torchvision.models.detection import maskrcnn_resnet50_fpn, \
    MaskRCNN_ResNet50_FPN_Weights
if os.path.exists('dog-cycle-car.png'):
    w = MaskRCNN_ResNet50_FPN_Weights.DEFAULT
    names = w.meta['categories'] # 91 COCO names, index 0 = \
    \background
    mrcnn = maskrcnn_resnet50_fpn(weights=w).eval().to(device)
    img = cv2.cvtColor(cv2.imread('dog-cycle-car.png'), cv2.COLOR_BGR2RGB)
    with torch.no_grad():
        out = mrcnn([TF.to_tensor(img).to(device)])[0]
        keep = out['scores'].cpu().numpy() >= 0.7
        boxes = out['boxes'].cpu().numpy()[keep]; labels = out['labels'].cpu().\
        \numpy()[keep]
```

```

    masks = out['masks'].cpu().numpy()[keep]
    rng = np.random.RandomState(42); COLORS = rng.randint(0, 255, (91, 3))
    fig, axes = plt.subplots(1, 3, figsize=(18, 6))
    axes[0].imshow(img); axes[0].set_title('Original', fontweight='bold');
    ↪axes[0].axis('off')
    axes[1].imshow(img)
    for b, l in zip(boxes, labels):
        axes[1].add_patch(plt.Rectangle((b[0], b[1]), b[2]-b[0], b[3]-b[1],
    ↪fill=False, edgecolor=COLORS[l]/255, lw=2.5))
        axes[1].text(b[0], b[1]-4, names[l], color='white', fontsize=9,
                     bbox=dict(facecolor=COLORS[l]/255, alpha=0.8, pad=1,
    ↪edgecolor='none'))
    axes[1].set_title('Bounding boxes (Faster R-CNN head)', fontweight='bold');
    ↪axes[1].axis('off')
    overlay = img.copy().astype(np.float32)
    for m, l in zip(masks, labels):
        overlay[m[0] > 0.5] = overlay[m[0] > 0.5] * 0.5 + COLORS[l] * 0.5
    axes[2].imshow(overlay.astype(np.uint8))
    axes[2].set_title('Instance masks (Mask R-CNN head)', fontweight='bold');
    ↪axes[2].axis('off')
    fig.suptitle('Mask R-CNN = Faster R-CNN + mask head (pretrained COCO)',
    ↪fontsize=13, fontweight='bold')
    plt.show(); del mrcnn; torch.cuda.empty_cache()
else:
    print('dog-cycle-car.png not found - skipping Mask R-CNN demo (fetch via
    ↪script.ps1 step 2).')

```

Mask R-CNN = Faster R-CNN + mask head (pretrained COCO)



1.4 3. Data + model instances

Build the Oxford-Pet loaders (downloaded on first use) and instantiate both variants — note they share encoder weights and differ only by `use_skip` (and the decoder input channels that skips add).

```
[8]: train_loader, test_loader, train_data, test_data = get_pet_loaders(
      'data', size=IMG_SIZE, batch_size=BATCH_SIZE, num_workers=NUM_WORKERS,
      ↪download=True)
print(f'train={len(train_data)} test={len(test_data)}')
for name in ['unet_resnet18', 'unet_resnet18_no_skip']:
    mdl = build_segmentation_model(name, pretrained=False)
    p = sum(x.numel() for x in mdl.parameters())
    print(f'{name:24s} use_skip={mdl.use_skip!s:5s} params={p:,}')

```

```
train=3680 test=3669
unet_resnet18          use_skip=True  params=37,572,099
unet_resnet18_no_skip  use_skip=False params=34,420,227


```

1.4.1 Dataset samples — image / mask / overlay

```
[9]: fig, axes = plt.subplots(3, 3, figsize=(10, 10))
for i in range(3):
    img, mask = train_data[i*100]
    axes[i][0].imshow(denormalize(img)); axes[i][0].set_title('Image');
    ↪axes[i][0].axis('off')
    axes[i][1].imshow(CLASS_COLORS[mask.numpy()]); axes[i][1].set_title('Mask');
    ↪axes[i][1].axis('off')
    axes[i][2].imshow(denormalize(img)); axes[i][2].imshow(CLASS_COLORS[mask.
    ↪numpy()], alpha=0.5)
    axes[i][2].set_title('Overlay'); axes[i][2].axis('off')
handles = [plt.Rectangle((0,0),1,1,color=CLASS_COLORS[i]/255) for i in range(3)]
fig.legend(handles, CLASS_NAMES, loc='lower center', ncol=3)
fig.suptitle('Oxford-IIIT Pet - image + segmentation mask', fontsize=13); plt.
    ↪show()

```

Oxford-IIIT Pet — image + segmentation mask



1.5 4. Train both variants (20 epochs each)

`seg_load_or_train` reuses a cached checkpoint+history if present, otherwise trains via the inlined `train_segmentation` — identical to `run.py --model <name> --dataset oxford_pet --epochs 20 --train`.

```
[10]: def seg_load_or_train(name, epochs):
    hist_p, ckpt_p = f'checkpoints/history_{name}.json', f'checkpoints/{name}.
    ↪pt'
    if not RETRAIN and os.path.exists(hist_p) and os.path.exists(ckpt_p):
        model = build_segmentation_model(name).to(device)
        model.load_state_dict(torch.load(ckpt_p, map_location=device))
        history = json.load(open(hist_p))
        print(f'[{name}] reusing cached checkpoint + history_
    ↪({len(history["miou"])} epochs)')
    else:
        set_seed(SEED)    # reseed right before build+train -> matches run.py
        model = build_segmentation_model(name).to(device)
        history = train_segmentation(model, train_loader, test_loader, device,
                                     epochs=epochs, lr=LR,
    ↪ckpt_path=f'{name}_pet.pt', tag=name)
        model.eval()
        return model, history

model_skip, hs = seg_load_or_train('unet_resnet18', SEG_EPOCHS)
model_noskip, hn = seg_load_or_train('unet_resnet18_no_skip', SEG_EPOCHS)
```

```
[unet_resnet18] epoch 1/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 01/20 | Loss 0.4123 | mIoU 0.6789 | 20.7s
[unet_resnet18] epoch 2/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 02/20 | Loss 0.2999 | mIoU 0.7051 | 18.9s
[unet_resnet18] epoch 3/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 03/20 | Loss 0.2596 | mIoU 0.7167 | 18.8s
[unet_resnet18] epoch 4/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 04/20 | Loss 0.2399 | mIoU 0.7099 | 18.9s
[unet_resnet18] epoch 5/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 05/20 | Loss 0.2152 | mIoU 0.7031 | 18.9s
[unet_resnet18] epoch 6/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 06/20 | Loss 0.1853 | mIoU 0.7456 | 18.9s
[unet_resnet18] epoch 7/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 07/20 | Loss 0.1644 | mIoU 0.7512 | 18.9s
[unet_resnet18] epoch 8/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 08/20 | Loss 0.1526 | mIoU 0.7476 | 18.8s
[unet_resnet18] epoch 9/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 09/20 | Loss 0.1496 | mIoU 0.7468 | 18.9s
```

```

[unet_resnet18] epoch 10/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 10/20 | Loss 0.1410 | mIoU 0.7473 | 18.9s
[unet_resnet18] epoch 11/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 11/20 | Loss 0.1312 | mIoU 0.7593 | 18.6s
[unet_resnet18] epoch 12/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 12/20 | Loss 0.1229 | mIoU 0.7598 | 18.8s
[unet_resnet18] epoch 13/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 13/20 | Loss 0.1181 | mIoU 0.7588 | 19.1s
[unet_resnet18] epoch 14/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 14/20 | Loss 0.1163 | mIoU 0.7499 | 19.0s
[unet_resnet18] epoch 15/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 15/20 | Loss 0.1152 | mIoU 0.7552 | 18.5s
[unet_resnet18] epoch 16/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 16/20 | Loss 0.1075 | mIoU 0.7539 | 18.8s
[unet_resnet18] epoch 17/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 17/20 | Loss 0.1035 | mIoU 0.7536 | 18.8s
[unet_resnet18] epoch 18/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 18/20 | Loss 0.1008 | mIoU 0.7527 | 18.8s
[unet_resnet18] epoch 19/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 19/20 | Loss 0.0985 | mIoU 0.7505 | 18.9s
[unet_resnet18] epoch 20/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 20/20 | Loss 0.0965 | mIoU 0.7504 | 18.8s
Saved → unet_resnet18_pet.pt  and  checkpoints/history_unet_resnet18.json
[unet_resnet18_no_skip] epoch 1/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 01/20 | Loss 0.6785 | mIoU 0.4601 | 16.5s
[unet_resnet18_no_skip] epoch 2/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 02/20 | Loss 0.4858 | mIoU 0.5053 | 16.1s
[unet_resnet18_no_skip] epoch 3/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 03/20 | Loss 0.4324 | mIoU 0.5331 | 16.5s
[unet_resnet18_no_skip] epoch 4/20:  0%|          | 0/230 [00:00<?, ?it/s]
Epoch 04/20 | Loss 0.3884 | mIoU 0.5688 | 16.5s
[unet_resnet18_no_skip] epoch 5/20:  0%|          | 0/230 [00:00<?, ?it/s]

```

```

Epoch 05/20 | Loss 0.3577 | mIoU 0.5938 | 16.5s
[unet_resnet18_no_skip] epoch 6/20: 0%|          | 0/230 [00:00<?, ?it/s]
Epoch 06/20 | Loss 0.3073 | mIoU 0.6447 | 16.2s
[unet_resnet18_no_skip] epoch 7/20: 0%|          | 0/230 [00:00<?, ?it/s]
Epoch 07/20 | Loss 0.2768 | mIoU 0.6542 | 16.1s
[unet_resnet18_no_skip] epoch 8/20: 0%|          | 0/230 [00:00<?, ?it/s]
Epoch 08/20 | Loss 0.2594 | mIoU 0.6556 | 16.5s
[unet_resnet18_no_skip] epoch 9/20: 0%|          | 0/230 [00:00<?, ?it/s]
Epoch 09/20 | Loss 0.2390 | mIoU 0.6564 | 16.3s
[unet_resnet18_no_skip] epoch 10/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 10/20 | Loss 0.2243 | mIoU 0.6627 | 16.3s
[unet_resnet18_no_skip] epoch 11/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 11/20 | Loss 0.2018 | mIoU 0.6777 | 16.3s
[unet_resnet18_no_skip] epoch 12/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 12/20 | Loss 0.1880 | mIoU 0.6805 | 16.6s
[unet_resnet18_no_skip] epoch 13/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 13/20 | Loss 0.1806 | mIoU 0.6818 | 16.5s
[unet_resnet18_no_skip] epoch 14/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 14/20 | Loss 0.1756 | mIoU 0.6820 | 16.2s
[unet_resnet18_no_skip] epoch 15/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 15/20 | Loss 0.1701 | mIoU 0.6841 | 16.4s
[unet_resnet18_no_skip] epoch 16/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 16/20 | Loss 0.1597 | mIoU 0.6851 | 16.1s
[unet_resnet18_no_skip] epoch 17/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 17/20 | Loss 0.1533 | mIoU 0.6857 | 16.0s
[unet_resnet18_no_skip] epoch 18/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 18/20 | Loss 0.1491 | mIoU 0.6844 | 16.5s
[unet_resnet18_no_skip] epoch 19/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 19/20 | Loss 0.1468 | mIoU 0.6876 | 16.3s
[unet_resnet18_no_skip] epoch 20/20: 0%|         | 0/230 [00:00<?, ?it/s]
Epoch 20/20 | Loss 0.1442 | mIoU 0.6873 | 16.4s
Saved → unet_resnet18_no_skip_pet.pt and
checkpoints/history_unet_resnet18_no_skip.json

```

1.6 5. Results — skip vs no-skip

1.6.1 Summary table (Exercise 1b)

```
[11]: tbl = pd.DataFrame({
    'Model': ['U-Net + ResNet-18 (skip)', 'U-Net + ResNet-18 (no skip)'],
    'Skip connections': ['Yes', 'No'],
    'Val mIoU': [round(hs['miou'][-1],4), round(hn['miou'][-1],4)], #
    ↪final-epoch = saved checkpoint = run.py --evaluate
    'Time/epoch (s)': [round(np.mean(hs['time']),1), round(np.
    ↪mean(hn['time']),1)]]
display(tbl)
print(f"skip improves Val mIoU by {hs['miou'][-1]-hn['miou'][-1]:+.4f}")
# Persist final-epoch mIoU to results/metrics.json (same keys/semantics as run.
    ↪py --evaluate)
record_metric('UNET_resnet18_miou', round(hs['miou'][-1], 4))
record_metric('UNET_resnet18_no_skip_miou', round(hn['miou'][-1], 4))
print('wrote results/metrics.json:', json.load(open('results/metrics.json')))
```


	Model	Skip connections	Val mIoU	Time/epoch (s)
0	U-Net + ResNet-18 (skip)	Yes	0.7504	18.9
1	U-Net + ResNet-18 (no skip)	No	0.6873	16.3

skip improves Val mIoU by +0.0631

```
wrote results/metrics.json: {'yolov3_pretrained_map50': 0.6017,
'yolov3_pretrained_eval_images': 1000, 'yolov4_pretrained_map50': 0.6927,
'yolov4_pretrained_eval_images': 1000, 'UNET_resnet18_miou': 0.7504,
'UNET_resnet18_no_skip_miou': 0.6873, 'yolov4_pretrained_map5095': 0.4723,
'yolov3_pretrained_map5095': 0.3505, 'yolov4_mse_trained_map50': 0.0004,
'yolov4_mse_trained_map5095': 0.0001, 'yolov4_ciou_trained_map50': 0.0008,
'yolov4_ciou_trained_map5095': 0.0003}
```


1.6.2 Training curves

```
[12]: e = range(1, len(hs['loss'])+1)
fig, ax = plt.subplots(1, 2, figsize=(13, 5))
ax[0].plot(e, hs['loss'], marker='o', label='skip', color='steelblue')
ax[0].plot(e, hn['loss'], marker='s', label='no-skip', color='firebrick')
ax[0].set_title('Training loss'); ax[0].set_xlabel('epoch'); ax[0].legend();
    ↪ax[0].grid(alpha=0.3)
ax[1].plot(e, hs['miou'], marker='o', label='skip', color='steelblue')
ax[1].plot(e, hn['miou'], marker='s', label='no-skip', color='firebrick')
ax[1].set_title('Validation mIoU'); ax[1].set_xlabel('epoch'); ax[1].legend();
    ↪ax[1].grid(alpha=0.3)
fig.suptitle('U-Net / ResNet-18 - skip vs no-skip', fontsize=13,
    ↪fontweight='bold'); plt.show()
```

1.6.3 Per-class IoU (final epoch) & best mIoU

```
[13]: pcs, pcn = hs['per_class'][-1], hn['per_class'][-1]
x = np.arange(3); w = 0.35
fig, ax = plt.subplots(1, 2, figsize=(13, 5))
ax[0].bar(x - w/2, pcs, w, label='skip', color='steelblue')
ax[0].bar(x + w/2, pcn, w, label='no-skip', color='firebrick')
ax[0].set_xticks(x); ax[0].set_xticklabels(CLASS_NAMES); ax[0].set_ylabel('IoU')
ax[0].set_title('Per-class IoU (final epoch)'); ax[0].legend(); ax[0].
    grid(axis='y', alpha=0.3)
ax[1].bar(['skip', 'no-skip'], [hs['miou'][-1], hn['miou'][-1]],
    color=['steelblue', 'firebrick'])
ax[1].set_title('Val mIoU (final epoch)'); ax[1].set_ylabel('mIoU'); ax[1].
    grid(axis='y', alpha=0.3)
for i, v in enumerate([hs['miou'][-1], hn['miou'][-1]]):
    ax[1].text(i, v + 0.005, f'{v:.3f}', ha='center', fontweight='bold')
fig.suptitle('Skip connections - per-class IoU & mIoU', fontsize=13,
    fontweight='bold'); plt.show()
```




1.6.4 Per-model prediction grids — Input / Ground Truth / Prediction / Overlay

```
[14]: def show_seg_predictions(model, name, n=5, step=50):
    fig, axes = plt.subplots(n, 4, figsize=(14, 3.6*n))
    for ax, t in zip(axes[0], ['Input', 'Ground Truth', 'Prediction', '
    ↪'Overlay']):
        ax.set_title(t, fontweight='bold')
    for r in range(n):
        img, mask = test_data[r*step]
        with torch.no_grad():
            pred = model(img.unsqueeze(0).to(device)).argmax(1).squeeze().cpu().
            ↪numpy()
            axes[r][0].imshow(denormalize(img))
            axes[r][1].imshow(CLASS_COLORS[mask.numpy()])
            axes[r][2].imshow(CLASS_COLORS[pred])
            axes[r][3].imshow(denormalize(img)); axes[r][3].
            ↪imshow(CLASS_COLORS[pred], alpha=0.5)
            for ax in axes[r]:
                ax.axis('off')
            handles = [plt.Rectangle((0,0),1,1,color=CLASS_COLORS[i]/255) for i in
            ↪range(3)]
            fig.legend(handles, CLASS_NAMES, loc='lower center', ncol=3)
            fig.suptitle(f'U-Net predictions - {name}', fontsize=13); plt.show()


show_seg_predictions(model_skip, 'UNET_resnet18')
show_seg_predictions(model_noskip, 'UNET_resnet18_no_skip')
```

U-Net predictions — unet_resnet18



 Pet  Background  Border


U-Net predictions — unet_resnet18_no_skip



1.6.5 Predictions side by side (same images) — note how no-skip blurs boundaries

```
[15]: rows = 4
fig, axes = plt.subplots(rows, 4, figsize=(14, 4*rows))
for ax, t in zip(axes[0], ['Input', 'Ground Truth', 'Skip', 'No-skip']):
    ax.set_title(t, fontweight='bold')
for r in range(rows):
    img, mask = test_data[r*60 + 20]
    with torch.no_grad():
        ps = model_skip(img.unsqueeze(0).to(device)).argmax(1).squeeze().cpu().
        ↪numpy()
        pn = model_noskip(img.unsqueeze(0).to(device)).argmax(1).squeeze().
        ↪cpu().numpy()
    axes[r][0].imshow(denormalize(img))
    axes[r][1].imshow(CLASS_COLORS[mask.numpy()])
    axes[r][2].imshow(CLASS_COLORS[ps])
    axes[r][3].imshow(CLASS_COLORS[pn])
    for ax in axes[r]:
        ax.axis('off')
fig.suptitle('Skip vs no-skip predictions - note boundary sharpness',
        ↪fontsize=13, fontweight='bold'); plt.show()
```

Skip vs no-skip predictions — note boundary sharpness



1.7 6. Exercise answers

1.7.1 (c) Why do skip connections help segmentation more than classification?

Classification only needs *what* — a single global label — so the network deliberately discards spatial detail: pooling/striding build compact, translation-invariant semantics and the final answer doesn't depend on *where* an edge sat. Segmentation needs *what and exactly where* — a label per pixel — so the decoder must rebuild sharp boundaries at full resolution, which is impossible from the coarse bottleneck alone (a blurry thumbnail). Skip connections copy the encoder's high-resolution feature maps (precise edges/texture/location) straight into the decoder, fusing deep semantics with shallow geometry. For classification those shallow features barely affect the pooled prediction, so skips add little — they are far more valuable for dense prediction.

1.7.2 (d) Which skip level hurts most when removed — first (64ch, highest res) or last (512ch, lowest res)?

The **first** skip (64-channel, $H/2$, highest resolution). It is the only path carrying full-resolution spatial detail into the final decoder stage; without it boundaries become blurry/blocky and thin structures (ear tips, legs, fur edges) are lost — precisely where pet IoU is decided. The last skip (512-channel, $H/32$) carries deep semantics that the adjacent bottleneck *already* encodes, so removing it is largely redundant.

1.8 7. Discussion

Skip connections improved Val mIoU from **0.687** (no-skip) to **0.750** (skip) — a **+0.063** gain for an otherwise identical encoder and decoder — and the side-by-side panels show the gain is concentrated at object **boundaries** (the Border/Pet classes), where the no-skip decoder must reconstruct fine edges from a coarse bottleneck and produces rounded, leaky masks. **Choose U-Net** when you need dense per-pixel labels for one or a few classes and don't need to separate instances — medical organ/tumour masks, road/sky 'stuff' — especially with limited data and a pretrained encoder. **Choose Mask R-CNN** when you must detect, separate and count individual object **instances** (this cat vs that cat) and also want boxes and labels, accepting its heavier two-stage cost.

1.9 8. Summary

- Two U-Net decoders share an **identical** ImageNet ResNet-18 encoder; the only difference is whether each decoder stage concatenates the matching encoder feature map (skip connections).
- Trained both on Oxford-IIIT Pet for 20 epochs with the same seed and schedule — skip connections raised Val mIoU (table above), with the gain concentrated at object **boundaries**.
- The **highest-resolution (first)** skip matters most; removing the lowest-resolution (last) skip is largely redundant because the bottleneck already encodes that semantics.
- Skip connections help **dense prediction** far more than classification, because segmentation must rebuild precise per-pixel boundaries that pooling/striding discard.